

# CSci 4061: Introduction to Operating Systems

Programming project 2

due: Tuesday March 12th, 2019

---

**Ground Rules.** You may choose to complete this project in a group of up to three students. Each group should turn in **one** copy with the names of all group members on it. The code must be originally written by your group. No code from outside the course texts and slides may be used—your code cannot be copied or derived from the Web, from past offerings, other students, programmer friends, etc. All submissions must compile and run on any CSE Labs machine located in KH 4-250. A zip file should be submitted through Canvas by 11:59pm on Tuesday, March 12th. **Note:** Do not publicize this assignment or your answer to the Internet, e.g., public GitHub repo.

**Objectives:** The main focus of this project is to implement (1) three Unix file system-related commands (`cd`, `du`, and `grep`) and (2) a unix shell which is capable of executing any Unix commands. Moreover, the shell should be able to handle redirection and pipes in addition to executing the commands.

## 1 Part 1: Implementation of Unix Commands

You have to implement the `du` and `grep` commands in their own C files which are already provided in the codebase. Those C files have its own *main* function. After you finish the implementation of the commands in their C files, the next step is to generate the executable files. To create their executable files, you can just execute ‘`make build`’ which builds the executable files for all the commands; the Makefile has been provided to you. Once you have the executable files for the commands, you should be able to run the commands from the terminal. For example, if you want to run the “`du`” command, you can just execute ‘`./du`’ like any other C program.

You have to implement all the below commands using basic system calls instead of just directly reusing (e.g., using *fork* and *exec*) the existing commands provided by the operating system. In the following, we describe the details about the commands. Note that all paths used in the following can be both absolute or relative.

### 1.1 `cd`

The `cd` command is used to change the working directory to the given path. “`..`” indicates the parent directory. Refer to the `cd` man page for more details.

**List of test cases:**

```
./cd <path>
./cd ..
```

(Hint: Use the `chdir` system call)

## 1.2 du

The `du` command gives the total sum of the size of the files/directories in the given path including all the files/directories in the sub-directories. You have to implement a modified version of the original `du -b` command. Given a path, you must print only one single value which is the total size (in bytes) taken by that path (including all sum of the size of all the files/directories in the sub-directories) in bytes.

### List of test cases:

```
./du <path>
```

If the path is not given, use the current working directory as the path.

(Hint: Use the `stat` system call and directory-related system calls. Recursion might be useful.)

Note: The output from your implementation will not be equal to `du -b`. You can verify whether your output value is correct by summing up all the size values displayed against each file when you run `ls -l` command on the same path.

## 1.3 grep

The `grep` command finds the given keyword from the given file and prints the line in which the keyword is present. Refer to `grep` man page for more details. You have to implement the regular `grep` command which takes in a file and a keyword and prints the line in which the keyword is present. And also, you have to implement the version of the `grep` command in which if the file name is not given, it takes the input text from the standard input and searches the keyword in inputted text.

### List of test cases:

```
./grep <keyword> <filename>
./grep <keyword>
```

In the second case, it should read the text from the standard input.

## 2 Part 2: Implementation of the Unix Shell

The second part of the project is to implement a unix shell that supports file redirection and pipes. The implementation should be done in the file '`shell.c`' in the given codebase. The shell should always display the current working directory along with a dollar while prompting the user to enter a command, as the original unix shell does. To differentiate from the original shell, add the tag `[4061-shell]` at the start. For example, if '`home/csci4061/project`' is the current working directory, it should display:

```
[4061-shell]/home/csci4061/project $
```

You must use `STDIN` to read the user input using the `read` system call but not `scanf`-based functions. And the shell should be able to parse any of the above commands and then execute the correct command along with user provided arguments. In the shell, you must use `fork` and `exec` to execute the command. If a specified command is not one of the above 3 commands implemented by you, you should call the system's original command executable directly using the `exec` function. If any error occurs during the execution of the given command, you have to output `'Command error'`.

**Additionally**, you have to implement `exit` command to terminate the shell. You can perform all the necessary steps to quit the shell (like freeing memory etc.) when handling the `exit` command.

## 2.1 File redirection

The shell should be able to handle redirection of any command's output to a file. In other words, the shell should be able to handle `'>'` and `'>>'` operators. For example,

```
[4061-shell]/home/csci4061/project $ ls -l > out.txt
```

The above command will redirect the output of `ls -l` command to a file named `out.txt`. The `'>'` operator will truncate the old contents in the file if any. And the `'>>'` operator will append the new contents to the old contents of the file if any.

(Hint: Use `dup2` system call)

## 2.2 Pipes

The shell should also be able to handle pipes between multiple commands. For simplicity, your shell should have the ability to handle one pipe between two commands. For example,

```
[4061-shell]/home/csci4061/project $ cat data.txt | grep student
```

The above command will search for the word `'student'` in the output of `'cat data.txt'`, which is the content of the file `'data.txt'`. Basically it will pipe the output of the `cat` command into the input of the `grep` command.

(Hint: Use pipe along with `fork` and `exec` system calls)

And also, the shell should support the combination of redirection and pipes in a single command. For example,

```
[4061-shell]/home/csci4061/project $ cat data.txt | grep student > out.txt
```

If you are curious, try implementing the support for multiple pipes. (But no extra credit will be provided)

## 3 Possible division of labour

- First person: Implementing `cd` and `exit` + Command parsing and execution logic
- Second person: Implementing `du` + file redirection
- Third person: Implementing `grep` + pipes

It is not necessary to follow the above division of labour. Feel free to change it as needed.  
(*Note: Start the project early. Don't wait for the midterm to be complete.*)

## 4 General notes

1. Use cselab machine (or similar environment) to implement the solution. Mac/Windows environment will not be suitable for this project as you have to read input from `STDIN`.
2. You are given a solution executable named `'sol'` against which you can compare the results from your implementation. (The solution might not run in Mac/Windows environment.).
3. You are given some of the helper functions in `'util.c'` which might be helpful. Feel free to change the implementation in `'util.c'` as needed. Alternatively, you can also modify the parser that you built for project 1.
4. The list of test cases given are the only test cases upon which your command will be tested. If your command works correctly for those basic test cases, you will get full credit for those commands.
5. Handle errors from system calls.

## 5 Deliverables

Students should upload to Canvas a zip file containing their C code, a makefile, and a README that includes the group member names, what each member contributed, any known bugs, test cases used (we will also use our own), whether the extra credit has been attempted, and any special instructions for running the code.

## 6 Grading Rubric

- 5% For correct README contents
- 5% Code quality such as using descriptive variable names and comments
- 40% Implementation of all the commands (`cd` and `exit` - 5% each, `du` and `grep` - 15% each)
- 25% Implementation of basic unix shell
- 10% Implementation of file redirection in shell
- 10% Implementation of pipes in shell
- 5% Ability to handle both redirection and pipe in a single command